

MACHINE LEARNING

March 20, 2025

Name : KARAMAGI HENRY ATUHAIRE

Registration Number : 2023/U/MMU/BCS/00523

1 Initial Model

```
[158]: # Importing the libraries

import pandas as pd

import seaborn as sns
import matplotlib.pyplot as plt
from scipy.sparse import csr_matrix
from joblib import parallel_backend

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
```

1. Loading the Iris dataset

```
[159]: # loading the dataset
df = pd.read_csv('iris.csv')
df.columns = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width', 'species']
```

```
[160]: # Features
X = df.iloc[:, :-1].values

# Target
y = df.iloc[:, -1].values
```

2. Preprocessing the Dataset : Scaling the Features

```
[161]: # Scaling the features
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

3. Split the Dataset into Training and Testing Sets

```
[162]: # Splitting data
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.
↳3, random_state=42)
```

4. Training a Logistic Regression Classifier on the Training Set

```
[163]: # The original model
original_model = LogisticRegression(max_iter=200)
original_model.fit(X_train, y_train)
```

```
[163]: LogisticRegression(max_iter=200)
```

5. Model Evaluation

```
[164]: # Model Evaluation
# make prediction
y_pred_original = original_model.predict(X_test)
accuracy_original = accuracy_score(y_test, y_pred_original)

# Calculating evaluation metrics for the original model
accuracy = accuracy_score(y_test, y_pred_original)
precision = precision_score(y_test, y_pred_original, average='weighted')
recall = recall_score(y_test, y_pred_original, average='weighted')
f1 = f1_score(y_test, y_pred_original, average='weighted')

# Results
print("Original Model Metrics")
print(f"Accuracy : {accuracy:.2f}")
print(f"Precision : {precision:.2f}")
print(f"Recall : {recall:.2f}")
print(f"F1 Score : {f1:.2f}")
```

Original Model Metrics

```
Accuracy : 0.91
Precision : 0.92
Recall : 0.91
F1 Score : 0.91
```

2 Cross Validation of the Original model

```
[165]: cv_scores = cross_val_score(original_model, X_scaled,y, cv=5)
print("Cross Validation with the Original Model")
print(f"C-V Scores: {cv_scores}")
print(f"Average C-V Accuracy: {cv_scores.mean():.2f}")
```

Cross Validation with the Original Model

C-V Scores: [0.96666667 1. 0.93333333 0.9 1.]

Average C-V Accuracy: 0.96

3 Optimization

3.0.1 Option 1: Parallel Processing

```
[166]: with parallel_backend('threading', n_jobs=-1): # using all available CPU cores
# Enabling parallel processing
parallel_model = LogisticRegression(max_iter=200, n_jobs=-1)

# fitting the data
parallel_model.fit(X_train, y_train)

# Calculating evaluation metrics for the parallel processing model
accuracy_parallel = accuracy_score(y_test, y_pred_parallel)
precision_parallel = precision_score(y_test, y_pred_parallel, average='weighted')
recall_parallel = recall_score(y_test, y_pred_parallel, average='weighted')
f1_parallel = f1_score(y_test, y_pred_parallel, average='weighted')
conf_matrix_parallel = confusion_matrix(y_test, y_pred_parallel)

# Results
print("\nParallel Processing Model Metrics")
print(f"Accuracy : {accuracy_parallel:.2f}")
print(f"Precision : {precision_parallel:.2f}")
print(f"Recall : {recall_parallel:.2f}")
print(f"F1 Score : {f1_parallel:.2f}")
```

Parallel Processing Model Metrics

Accuracy : 1.00

Precision : 1.00

Recall : 1.00

F1 Score : 1.00

3.0.2 Option 2: Efficient Algorithm using Solver

```
[167]: # Using the solver parameter
efficient_model = LogisticRegression(solver='saga', max_iter=200)
efficient_model.fit(X_train, y_train)
```

```

# Calculate evaluation metrics for the efficient algorithm model
accuracy_efficient = accuracy_score(y_test, y_pred_efficient)
precision_efficient = precision_score(y_test, y_pred_efficient,
    ↪average='weighted')
recall_efficient = recall_score(y_test, y_pred_efficient, average='weighted')
f1_efficient = f1_score(y_test, y_pred_efficient, average='weighted')
conf_matrix_efficient = confusion_matrix(y_test, y_pred_efficient)

# Print the results
print("\nSolver Model Metrics")
print(f"Accuracy    : {accuracy_efficient:.2f}")
print(f"Precision    : {precision_efficient:.2f}")
print(f"Recall       : {recall_efficient:.2f}")
print(f"F1 Score     : {f1_efficient:.2f}")

```

```

Solver Model Metrics
Accuracy    : 1.00
Precision   : 1.00
Recall      : 1.00
F1 Score    : 1.00

```

3.0.3 Option 3: Sparse Matrices

```

[168]: # Converting to a sparse matrix
X_sparse = csr_matrix(X_scaled)
X_train_sparse, X_test_sparse, y_train, y_test = train_test_split(X_sparse, y,
    ↪test_size=0.3, random_state=42)

sparse_model = LogisticRegression(max_iter=200)
sparse_model.fit(X_train_sparse, y_train)

# Calculate evaluation metrics for the sparse matrices model
accuracy_sparse = accuracy_score(y_test, y_pred_sparse)
precision_sparse = precision_score(y_test, y_pred_sparse, average='weighted')
recall_sparse = recall_score(y_test, y_pred_sparse, average='weighted')
f1_sparse = f1_score(y_test, y_pred_sparse, average='weighted')
conf_matrix_sparse = confusion_matrix(y_test, y_pred_sparse)

# Print the results
print("\nSparse Matrices Model Metrics")
print(f"Accuracy    : {accuracy_sparse:.2f}")
print(f"Precision    : {precision_sparse:.2f}")
print(f"Recall       : {recall_sparse:.2f}")
print(f"F1 Score     : {f1_sparse:.2f}")

```

```

Sparse Matrices Model Metrics

```

```
Accuracy   : 1.00
Precision  : 1.00
Recall     : 1.00
F1 Score   : 1.00
```

4 Comparison of the Performance of the models

```
[169]: print("\nModel Performance Comparison")
print(f"Original Model Accuracy : {accuracy_original:.2f}")
print(f"Parallel Model Accuracy : {accuracy_parallel:.2f}")
print(f"Solver Model Accuracy   : {accuracy_efficient:.2f}")
print(f"Sparse Model Accuracy   : {accuracy_sparse:.2f}")
```


```
Model Performance Comparison
Original Model Accuracy : 0.91
Parallel Model Accuracy : 1.00
Solver Model Accuracy   : 1.00
Sparse Model Accuracy   : 1.00
```

5 Visualization

```
[170]: # 1. selecting the model with the highest accuracy
best_model = parallel_model if accuracy_parallel >= max(accuracy_original,
    ↳ accuracy_efficient, accuracy_sparse) else \
    efficient_model if accuracy_efficient >= max(accuracy_original,
    ↳ accuracy_parallel, accuracy_sparse) else \
    sparse_model if accuracy_sparse >= max(accuracy_original,
    ↳ accuracy_parallel, accuracy_efficient) else \
    original_model

y_pred_best = best_model.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred_best)
```

```
[171]: # 2. Plotting the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='BuPu',
    xticklabels=np.unique(y), yticklabels=np.unique(y))
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Best Model Confusion Matrix')
plt.show()
```



6 Cross Validation of the Best model

```
[172]: cv_scores = cross_val_score(best_model, X_scaled, y, cv=5)
print("Cross Validation with the Best Model")
print(f"C-V Scores: {cv_scores}")
print(f"Average C-V Accuracy: {cv_scores.mean():.2f}")
```

Cross Validation with the Best Model

C-V Scores: [0.96666667 1. 0.93333333 0.9 1.]

Average C-V Accuracy: 0.96

```
[ ]:
```